

COMP3314 Assignment 3

Image Classification Report

Team name: 99%

Chen Yuxuan 3036128028

Li Yufei 3036128432

Liu Yantong 3036127012

April 19, 2025

Contents

1	Dataset Analysis	4
1.1	Load Data	4
1.2	Statistics on the number of categories	5
1.2.1	Dataset Size	5
1.2.2	Category Label Count	5
1.2.3	Other Statistics	6
1.3	Image Visualization	7
1.3.1	One Example for Each Category	7
1.3.2	Ten Examples for Each Category	7
2	Classifiers Comparison	9
2.1	Basic Workflow for testing classifiers	9
2.2	The performance of three classifiers	10
2.3	Analysis and Final choice	13
2.4	Final Choice	13
3	Final Model	14
3.1	Overall Model Structure	14
3.2	Feature Extraction	14
3.2.1	HOG (Histogram of oriented gradients)	14
3.2.2	EOG (Extract Edge Orientation Gradient)	16
3.2.3	ORB(Oriented FAST and Rotated BRIEF)	17
3.2.4	HIST(Histogram of Intensity and Surface Texture)	19
3.2.5	LBP (Local Binary Patterns)	20
3.2.6	SIFT(Scale-Invariant Feature Transform)	20
3.2.7	HSV (Hue-Saturation-Value)	22
3.3	Data Augmentation	22
3.3.1	Horizontal Flip	22
3.3.2	Random Crop	23
3.3.3	Random Rotation	23
3.3.4	Random Color Jittering	24
3.3.5	Approximate Final Size	25
3.4	StandardScaler	26
3.5	PCA	26
3.6	SVM Classifier	27

4	Final Result	27
4.1	Performance Analysis	27
4.2	Future Work	28

Abstract

1 Dataset Analysis

In this section, we present an overview of the dataset, analyzing the statistics on the number of categories. Additionally, we will visualize representative examples from each category to illustrate the diversity and characteristics inherent within the dataset. This analysis aims to provide a foundational understanding of the dataset's structure, which is essential for subsequent modeling and interpretation.

1.1 Load Data

We load the data and subsequently convert both the training and testing datasets into NumPy arrays, which are easier for our model to process.

```
csv_path, csv_test_path = "./data/train.csv", "./data/test.csv"
img_dir, img_dir_test = "./data/train_imgs", "./data/test_imgs"
data_train = pd.read_csv(csv_path)
data_test = pd.read_csv(csv_test_path)
X_train, y_train, X_test, y_test = [], [], [], []

for _, row in data_train.iterrows():
    img_path = os.path.join(img_dir, row.iloc[0])
    label = int(row.iloc[1])
    img = Image.open(img_path).convert("RGB")
    img = np.array(img).flatten()
    X_train.append(img)
    y_train.append(label)

for _, row in data_test.iterrows():
    img_path = os.path.join(img_dir_test, row.iloc[0])
    label = int(row.iloc[1])
    img = Image.open(img_path).convert("RGB")
    img = np.array(img).flatten()
    X_test.append(img)
    y_test.append(label)

X_train, y_train = np.array(X_train), np.array(y_train)
X_test, y_test = np.array(X_test), np.array(y_test)
X_train_flatten = X_train.reshape(X_train.shape[0], -1)
X_test_flatten = X_test.reshape(X_test.shape[0], -1)
```

Listing 1: Load Data

1.2 Statistics on the number of categories

1.2.1 Dataset Size

```
print("The shape of the training set: ", X_train.shape)
print("The shape of the testing set: ", X_test.shape)
print("Training set (flattened): ", X_train_flatten.shape)
print("Testing set (flattened): ", X_test_flatten.shape)
print("The size of the label of the training set: ", y_train.shape)
print("The size of the label of the testing set: ", y_test.shape)
```

Listing 2: Print the shape of the datasets

- **Size Statistics**
- The shape of the training set: (50000, 32, 32, 3)
- The shape of the testing set: (10000, 32, 32, 3)
- Training set (flattened): (50000, 3072)
- Testing set (flattened): (10000, 3072)
- The size of the label of the training set: (50000,)
- The size of the label of the test set: (10000,)

1.2.2 Category Label Count

```
# Count the labels
count_labels = np.unique(y_train, return_counts=True)

# Create DataFrame
label_counts_df = pd.DataFrame({
    "Label": count_labels[0],
    "Count": count_labels[1]
}).set_index("Label")

# Print DataFrame
print("Count of each label:")
print(label_counts_df)
```

Listing 3: Count the number of each category

Label 0	Label 1	Label 2	Label 3	Label 4
5038	5016	5032	4991	4982
Label 5	Label 6	Label 7	Label 8	Label 9
4967	4985	4998	5002	4989

Table 1: Label Counts in training set




Figure 1: Label Counts in training set

1.2.3 Other Statistics

```
print("Other Statistics for the training set label:")
print(pd.Series(y_train).describe())
```

Listing 4: Print other statistics of the labels of the training set

```
count    50000.00000
mean         4.49258
std         2.87539
min          0.00000
25%          2.00000
50%          4.00000
75%          7.00000
max          9.00000
```

1.3 Image Visualization

1.3.1 One Example for Each Category

```
def visualize_images_by_label(img_dir, data, y):
    unique_labels = np.unique(y)
    fig, axes = plt.subplots(2, 5, figsize=(15, 6))
    for i, label in enumerate(unique_labels[:10]):
        indices = np.where(y == label)[0]
        idx = np.random.choice(indices)
        img = Image.open(os.path.join(img_dir, data.iloc[idx, 0]))
        row = i // 5
        col = i % 5
        ax = axes[row, col]
        ax.imshow(img)
        ax.set_title(f"Label {label}")
        ax.axis('off')
    for j in range(i+1, 10):
        row = j // 5
        col = j % 5
        axes[row, col].axis('off')
    plt.tight_layout()
    plt.savefig("1 example.png", format='png', bbox_inches='tight')
    plt.show()
visualize_images_by_label(img_dir, data_train, y_train)
```



Figure 2: Visualization of one example for each category

1.3.2 Ten Examples for Each Category

```

count_labels = np.unique(y_train, return_counts=True)
label_counts_df = pd.DataFrame({
    "Label": count_labels[0],
    "Count": count_labels[1]
}).set_index("Label")

def get_example_images(label, num_examples=10):
    indices = np.where(y_train == label)[0]
    selected_indices = np.random.choice(indices,
        size=min(num_examples, len(indices)), replace=False)
    example_images = [os.path.join(img_dir, data_train.iloc[i, 0])
        for i in selected_indices]
    return " ".join([f"<img src='{img}' width='50' />"
        for img in example_images])

label_counts_df["Example Images"] =
label_counts_df.index.map(get_example_images)
HTML(label_counts_df.to_html(escape=False))

```

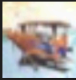
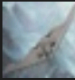





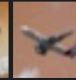






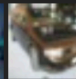





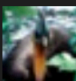
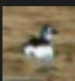
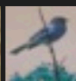
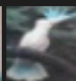
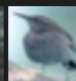

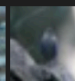
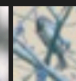
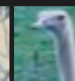
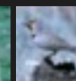
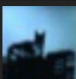

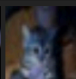
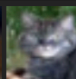
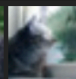
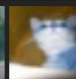
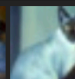
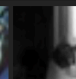
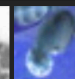
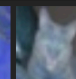
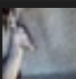
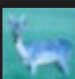
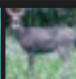
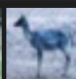
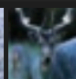
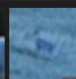
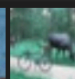
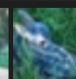
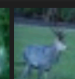
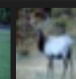
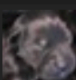
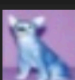
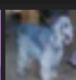
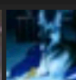
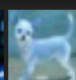
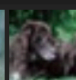
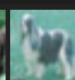
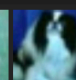
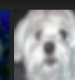
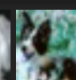
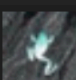
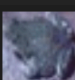
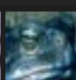
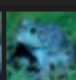
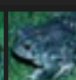
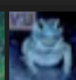
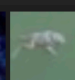
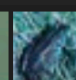
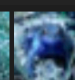
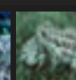
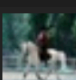
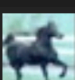
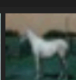
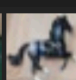
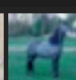
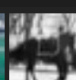
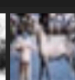
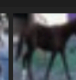
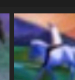
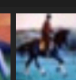
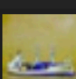
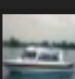
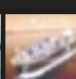
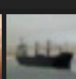
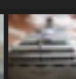

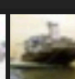
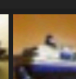
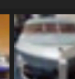
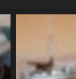
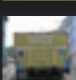
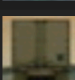
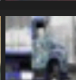
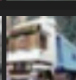
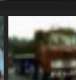
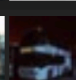
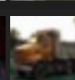
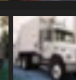
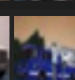
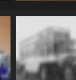
	Count	Example Images									
Label											
0	5038										
1	5016										
2	5032										
3	4991										
4	4982										
5	4967										
6	4985										
7	4998										
8	5002										
9	4989										

Figure 3: Visualization of ten example for each category

2 Classifiers Comparison

In this section, we'll explain how we tested and compared three different classifiers: KNN(K Nearest Neighbor), SVM(Support Vector Machine), RF(Random Forest) using the original dataset and the dataset with feature extraction(explained in Section 4).

2.1 Basic Workflow for testing classifiers

For the independence of variables, we use the following workflow for each models to ensure that factors such as the training set will not affect the three classifiers.

In this code, we use Pipeline to manage the process and the `classification_report` function of sklearn to estimate the performance, making the code more readable and satisfying the independence of variables.

```
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
#####
#build the pipeline
#####
XX_model = make_pipeline(
    StandardScaler(),
    PCA(n_components=0.75, random_state=SEED),
    XX())

#####
#Train and predict
#####
XX_model.fit(X_train_featured, y_train_aug)

y_pred_XX = XX_model.predict(X_val_featured)
accuracy_XX = accuracy_score(y_val, y_pred_XX)
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from matplotlib import pyplot as plt
```

```

#####
#Report
#####
# Classification Report
report = classification_report(y_val, y_pred_XX)

print("Classification Report:")
print(classification_report(y_val, y_pred_XX))

output_file = "classification_report.txt"

with open(output_file, "w") as f:
    f.write("Classification Report:\n")
    f.write(report)

print(f"Classification report saved to {output_file}")

```

2.2 The performance of three classifiers

The performance of three classifiers with origin data is shown below:

SVM Validation Accuracy: 0.4978

Classification Report:

	precision	recall	f1-score	support
0	0.53	0.55	0.54	1011
1	0.58	0.60	0.59	1051
2	0.38	0.38	0.38	985
3	0.35	0.33	0.34	983
4	0.41	0.41	0.41	968
5	0.44	0.39	0.41	975
6	0.49	0.59	0.54	1022
7	0.59	0.55	0.57	987
8	0.60	0.62	0.61	996
9	0.58	0.54	0.56	1022
accuracy			0.50	10000
macro avg	0.50	0.50	0.49	10000
weighted avg	0.50	0.50	0.50	10000

KNN Validation Accuracy: 0.4029

Classification Report:

	precision	recall	f1-score	support
0	0.41	0.57	0.48	1011
1	0.55	0.40	0.47	1051
2	0.28	0.38	0.32	985
3	0.31	0.21	0.25	983
4	0.28	0.38	0.32	968
5	0.42	0.26	0.32	975
6	0.36	0.51	0.42	1022
7	0.57	0.36	0.44	987
8	0.49	0.57	0.53	996
9	0.59	0.36	0.45	1022
accuracy			0.40	10000
macro avg	0.43	0.40	0.40	10000
weighted avg	0.43	0.40	0.40	10000

RF Validation Accuracy: 0.4573

Classification Report:

	precision	recall	f1-score	support
0	0.53	0.52	0.52	1011
1	0.50	0.56	0.53	1051
2	0.40	0.34	0.37	985
3	0.32	0.29	0.30	983
4	0.41	0.38	0.39	968
5	0.38	0.35	0.37	975
6	0.45	0.51	0.48	1022
7	0.53	0.47	0.50	987
8	0.53	0.62	0.57	996
9	0.48	0.52	0.50	1022
accuracy			0.46	10000
macro avg	0.45	0.46	0.45	10000
weighted avg	0.45	0.46	0.45	10000

After feature extractions(which will be detailed explained in the following sections), the performance is:

SVM Validation Accuracy: 0.7553

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.83	0.81	1011
1	0.87	0.87	0.87	1051

2	0.65	0.66	0.65	985
3	0.55	0.56	0.56	983
4	0.68	0.73	0.70	968
5	0.64	0.61	0.62	975
6	0.83	0.81	0.82	1022
7	0.84	0.80	0.82	987
8	0.85	0.85	0.85	996
9	0.85	0.83	0.84	1022
accuracy			0.76	10000
macro avg	0.75	0.75	0.75	10000
weighted avg	0.76	0.76	0.76	10000

KNN Validation Accuracy: 0.5914

Classification Report:

	precision	recall	f1-score	support
0	0.75	0.65	0.70	1011
1	0.77	0.81	0.79	1051
2	0.63	0.41	0.50	985
3	0.63	0.14	0.23	983
4	0.42	0.60	0.50	968
5	0.69	0.26	0.38	975
6	0.35	0.92	0.51	1022
7	0.84	0.62	0.72	987
8	0.72	0.75	0.73	996
9	0.74	0.70	0.72	1022
accuracy			0.59	10000
macro avg	0.65	0.59	0.58	10000
weighted avg	0.66	0.59	0.58	10000

RF Validation Accuracy: 0.5410

Classification Report:

	precision	recall	f1-score	support
0	0.59	0.63	0.61	1011
1	0.65	0.74	0.69	1051
2	0.45	0.38	0.41	985
3	0.36	0.26	0.30	983
4	0.42	0.40	0.41	968
5	0.42	0.49	0.45	975

	6	0.58	0.66	0.62	1022
	7	0.62	0.59	0.61	987
	8	0.61	0.65	0.63	996
	9	0.62	0.59	0.61	1022
	accuracy			0.54	10000
	macro avg	0.53	0.54	0.53	10000
	weighted avg	0.53	0.54	0.54	10000

2.3 Analysis and Final choice

Before Feature Extraction

SVM: Moderate accuracy (50%), balanced but mediocre performance.

KNN: Poor accuracy (40%), high variability across classes.

RF: Slightly better (46%), but still struggles with some classes.

After Feature Extraction

SVM: Best performance (75.5% accuracy), consistent across all classes.

KNN: Improved (59%), but suffers from extreme recall imbalances (e.g., some classes <20%).

RF: Minor improvement (54%), still lacks robustness in difficult classes.

2.4 Final Choice

We select SVM as our final classifier based on two key advantages:

- **Superior Classification Performance:**
 - Achieves highest accuracy (75.5%) among all candidates
 - Maintains stable performance across all classes (F1-score range: 0.56-0.87)
 - Demonstrates 25.7% absolute improvement after feature extraction
- **Dimensionality Robustness:**
 - Effectively handles high-dimensional feature space
 - Shows strongest synergy with our feature extraction pipeline
 - Margin-based optimization prevents overfitting

And since SVM excels with enhanced features, adding weaker models (KNN/RF) would add noise, not value. So we didn't ensemble other weaker models.

3 Final Model

3.1 Overall Model Structure

The overall structure of our classification model is as follows. We used data augmentation techniques such as horizontal flipping, random cropping, random rotation, and random color jittering to increase the size of the training dataset to approximately 2.3 times the original. Then, we employed the HOG method, supplemented by SIFT, HIST, EOG, and ORB methods, for feature extraction from the training set, while resizing the features during extraction for finer detail. Finally, We created a pipeline comprising StandardScaler, PCA, and SVM for training.

3.2 Feature Extraction

In the feature extraction process, we utilized several methods to extract features from the images. The following sections will detail the feature extraction methods we employed, including HOG, EOG, ORB, HIST, and LBP. We will also provide visualizations of the features extracted by each method.

3.2.1 HOG (Histogram of oriented gradients)

The Histogram of Oriented Gradients (HOG) is a feature descriptor widely used in computer vision for object detection. It quantifies local gradient orientations by dividing an image into cells, computing gradient magnitudes and orientations within each cell, and aggregating them into orientation-based histograms. These histograms are normalized within overlapping spatial blocks to enhance illumination invariance. HOG effectively captures edge structure and texture patterns, achieving robustness to geometric and photometric variations.




Figure 4: Visualization of HOG

In our implementation, we firstly resize the image to 128x128 pixels, then split the image into three channels (RGB), and finally apply HOG to each channel. The HOG features are concatenated into a single feature vector for each image. The following code demonstrates the process:

```
def HOG_extractor(image, target_size=(128, 128)):
    if len(image.shape) == 1:
        image = image.reshape(32, 32, 3) # Adjust this based on your original image

    resized_img = cv2.resize(image, target_size, interpolation=cv2.INTER_LINEAR)
    channels = cv2.split(resized_img)

    hog_kwargs = {
        "_winSize": (128, 128),
        "_blockSize": (64, 64),
        "_blockStride": (16, 16),
        "_cellSize": (16, 16),
        "_nbins": 10,
        "_derivAperture": 1,
        "_winSigma": -1,
        "_histogramNormType": 0,
        "_L2HysThreshold": 0.2,
        "_gammaCorrection": True,
        "_nlevels": 64,
        "_signedGradient": True
    }

    hog = cv2.HOGDescriptor(**hog_kwargs)
```

```

hog_features = []
for channel in channels:
    hog_feature = hog.compute(channel)
    hog_features.append(hog_feature.flatten())

return np.concatenate(hog_features)

```

After adopting the HOG method, the performance of the SVM rbf kernel model improved significantly, the average accuracy increased around 20%. As a result, we decided to use HOG as our main feature extraction method.

3.2.2 EOG (Extract Edge Orientation Gradient)

The EOG extraction method offers a computationally efficient approach for capturing essential structural information by combining Canny edge detection with gradient analysis, which enhances boundary sensitivity through precise thresholding and Sobel filtering while suppressing noise. Its grayscale conversion and feature vector flattening optimize processing efficiency and storage requirements, while image resizing maintains dimensional consistency across varied inputs. The method simultaneously preserves rich directional texture characteristics through gradient magnitude (quantifying edge strength) and orientation (encoding angular patterns), enabling comprehensive yet compact representation of both geometric contours and subtle surface variations in visual data.




Figure 5: Visualization of EOG

After add the EOG, the accuracy increased around 1% compared to the HOG method, from 0.7670 to 0.7792.

We try both 3 channel EOG and 1 channel EOG, however, the performance of 3 channel EOG is not as good as 1 channel EOG. The following code demonstrates the process. The accuracy of 3 channel EOG is 0.7787, while the accuracy of 1 channel EOG is around 0.7792. Therefore, we only use 1 channel EOG as our feature extraction method. We indicate the abnormal

result may be caused by overfitting, as the 3 channel EOG has more features than the 1 channel EOG.

The implementation of EOG is as follows:

```
def EOG_extractor(image, target_size=(32, 32)):
    if len(image.shape) == 1:
        image = image.reshape(32, 32, 3)

    resized_img = cv2.resize(image, target_size, interpolation=cv2.INTER_CUBIC)
    gray_img = cv2.cvtColor(resized_img, cv2.COLOR_RGB2GRAY)

    # Apply Canny edge detection
    edges = cv2.Canny(gray_img, threshold1=100, threshold2=200)
    # Compute gradients in x and y directions
    grad_x = cv2.Sobel(edges, cv2.CV_64F, 1, 0, ksize=3)
    grad_y = cv2.Sobel(edges, cv2.CV_64F, 0, 1, ksize=3)
    # Compute gradient magnitude and orientation
    magnitude = np.sqrt(grad_x**2 + grad_y**2)
    orientation = np.arctan2(grad_y, grad_x)
    # Flatten the magnitude and orientation into a feature vector
    eog_feature = np.concatenate([magnitude.flatten(), orientation.flatten()])

    return eog_feature
```

Meanwhile, the effect of EOG feature extraction on resized images is not satisfactory, which reduces the final accuracy by 1%, drop to 0.7655.

3.2.3 ORB(Oriented FAST and Rotated BRIEF)

ORB is a fast and efficient feature extraction algorithm used in computer vision. It combines the FAST keypoint detector, which identifies keypoints based on intensity changes, with the BRIEF descriptor, which encodes local image patches into binary strings for matching. ORB introduces improvements such as orientation estimation for rotation invariance and a more robust descriptor through rotated BRIEF.




Figure 6: Visualization of ORB

We try to use the ORB as our supplementary feature extraction method. The implementation of ORB is as follows:

```
def ORB_extractor(image, target_size=(256, 256), n_features=500, fixed_length=1000):
    orb = cv2.ORB_create(nfeatures=n_features)
    if len(image.shape) == 1:
        image = image.reshape(32, 32, 3)
    resized_img = cv2.resize(image, target_size, interpolation=cv2.INTER_CUBIC)
    gray_img = cv2.cvtColor(resized_img, cv2.COLOR_RGB2GRAY)

    # Detect keypoints and compute descriptors
    keypoints, descriptors = orb.detectAndCompute(gray_img, None)
    # If no descriptors are found, return a zero vector
    if descriptors is None:
        descriptors = np.zeros((1, 32), dtype=np.uint8)

    descriptors = descriptors.flatten()
    if len(descriptors) < fixed_length:
        descriptors = np.pad(descriptors, (0, fixed_length - len(descriptors)), mode='constant')
    elif len(descriptors) > fixed_length:
        descriptors = descriptors[:fixed_length]

    return descriptors.flatten()
```

However, the ORB feature extraction process has become a performance bottleneck, degrading the overall accuracy by around 1%.

3.2.4 HIST(Histogram of Intensity and Surface Texture)

The Histogram of Intensity and Surface Texture (HIST) feature extraction method quantifies spatial-intensity distributions and textural patterns by constructing multi-dimensional histograms across localized image regions. It employs adaptive binning to capture intensity gradients and co-occurrence statistics, coupled with spatial pooling to preserve structural context. Normalization techniques mitigate illumination variations while entropy-based weighting enhances discriminative power. This approach balances computational efficiency with robustness to affine transformations, making it particularly effective for texture classification, material recognition, and illumination-invariant object detection in computer vision application




Figure 7: Visualization of ORB

Our implementation is based on the OpenCV library, which provides a convenient function for histogram calculation. The following code demonstrates the process:

```
def HIST_extractor(image, target_size=(256, 256), bins=(8, 8, 8)):
    if len(image.shape) == 1:
        image = image.reshape(32, 32, 3)
    resized_img = cv2.resize(image, target_size, interpolation=cv2.INTER_LINEAR)

    # Compute the color histogram for each channel (B, G, R)
    hist = cv2.calcHist([resized_img], [0, 1, 2], None, bins, [0, 256, 0, 256, 0, 256])
    # Normalize the histogram
    hist = cv2.normalize(hist, hist).flatten()

    return hist
```

By adding the HIST method, the accuracy increased around 1% compared to the single EOG + HOG method, from 0.7670 to 0.7792.

3.2.5 LBP (Local Binary Patterns)

The Local Binary Patterns (LBP) method encodes texture information by comparing pixel intensities with their circular neighborhood.




Figure 8: Visualization of ORB

Our OpenCV-based implementation adopts uniform pattern mapping with histogram normalization:

```
def LBP_extractor(image, target_size=(64, 64)):
    if len(image.shape) == 1:
        image = image.reshape(32, 32, 3) # Adjust this based on your original image
    resized_img = cv2.resize(image, target_size, interpolation=cv2.INTER_LINEAR)

    # Convert to grayscale
    gray_img = rgb2gray(resized_img)
    # Compute LBP features
    lbp = local_binary_pattern(gray_img, P=8, R=1, method='default')
    # Flatten the LBP feature vector
    return lbp.flatten()
```

While theoretically effective for texture characterization, our implementation exhibited around 1.3% accuracy drop (from 0.7792 to 0.7668) when integrated into the feature ensemble (both default and uniform lbp), suggesting potential implementation challenges or feature conflicts.

3.2.6 SIFT(Scale-Invariant Feature Transform)

SIFT detects and describes distinctive image features invariant to scale, rotation, and illumination by: 1) Identifying keypoints via Difference-of-Gaussians extrema in scale-space pyramids; 2) Assigning orientation-based descriptors using gradient orientation histograms in localized regions. Its 128-dimensional

vectors enable robust matching for object recognition, 3D reconstruction, and image alignment despite viewpoint changes or partial occlusion.




Figure 9: Visualization of ORB

Our implementation is as follows:

```
def SIFT_extractor(image, n_features=128):
    sift = cv2.SIFT_create()
    if not isinstance(image, np.ndarray):
        image = np.array(image)
    if len(image.shape) == 1:
        image = image.reshape(32, 32, 3)

    channels = cv2.split(image)
    descriptors_list = []

    for channel in channels:
        keypoints, descriptors = sift.detectAndCompute(channel, None)
        # If no descriptors are found, return an empty array
        if descriptors is None:
            descriptors = np.zeros((1, n_features)).flatten()
        descriptors_mean = np.mean(descriptors, axis=0)
        descriptors_list.append(descriptors_mean)

    combined_descriptors = np.concatenate(descriptors_list)

    return combined_descriptors
```

However, adding the SIFT has no significant effect on the accuracy, which is around 0.7792. To save more memory for the training process, we decided

to remove the SIFT method from our final model.

3.2.7 HSV (Hue-Saturation-Value)

The Hue-Saturation-Value (HSV) color model is a key tool in computer vision for color-based analysis. It represents colors through three components: Hue (color type, 0° - 360°), Saturation (color purity, 0%-100%), and Value (brightness, 0%-100%). Unlike RGB, HSV separates color information from intensity, making it more robust to lighting changes. This separation allows for effective color segmentation and object detection while maintaining intuitive color manipulation. HSV's perceptual alignment with human vision and its illumination invariance make it particularly useful for applications like skin detection and traffic sign recognition.




Figure 10: Visualization of HSV

3.3 Data Augmentation

As we are restricted to using other data sets, we enhance the training set by processing the training images in the existing data set. In our final model, we applied horizontal flipping to each image, along with a 26×26 crop at the four corners and the center. For each image, there is a probability of 0.1 to perform a random local crop, with the crop size ranging from 24×24 to 32×32 and the crop position being random. Furthermore, there is a probability of 0.1 to apply a random image rotation, with a maximum angle of ± 18 degrees. There is also a probability of 0.1 that random color jittering will occur. If color jittering is applied, the brightness, contrast, and hue are randomly adjusted within the range of 0.8 to 1.2.

3.3.1 Horizontal Flip

Horizontal flipping is a simple yet effective augmentation technique. By mirroring the images along the vertical axis, we can create a new version of each

image. This is particularly useful in scenarios where the orientation of objects is not fixed (e.g., animals, cars, etc.). Since many visual patterns remain unchanged under horizontal flipping, this technique helps our model learn to recognize features regardless of their orientation.

```
flipped_img = pil_img.transpose(Image.FLIP_LEFT_RIGHT)
```

Listing 5: Horizontal Flips




Figure 11: [Left]Original [Right]Horizontal Flip

3.3.2 Random Crop

Random local cropping involves selecting a random portion of the image. With a probability of 0.1, we crop images to sizes between 24x24 and 32x32 pixels (recovered to 32x32 pixels at the end):

- **Crop Size:** By varying the crop size, we expose the model to different scales of the same object, which can help it become invariant to changes in scale.
- **Random Position:** The crop position is also random, ensuring that the model learns to identify relevant features regardless of their location within the image. This simulates the presence of objects at different positions and reinforces the model's ability to generalize.

3.3.3 Random Rotation

Random rotation adds another layer of variability by slightly rotating images within a maximum angle of ± 18 degrees. Many objects can appear at various angles in real-world scenarios. By allowing the model to see rotated images,

```
def random_crop(image, crop_size):
    width, height = image.size
    left = random.randint(0, width - crop_size)
    upper = random.randint(0, height - crop_size)
    right = left + crop_size
    lower = upper + crop_size
    return image.crop((left, upper, right, lower))
```

Listing 6: Random Crop




Figure 12: [Left]Original [Right]Crop

it learns to recognize patterns and features even when they are not perfectly aligned. The rotation is kept within ± 18 degrees to prevent excessive distortion, which could lead to loss of important features or misinterpretation of the image content.

```
def random_rotation(image, max_angle=18):
    angle = random.randint(-max_angle, max_angle)
    return image.rotate(angle)
```

Listing 7: Random Rotation

3.3.4 Random Color Jittering

Color jittering introduces variability in the color properties of the images. This technique is particularly useful in real-world scenarios where lighting conditions can vary significantly.

- **Brightness Adjustment:** Randomly altering the brightness helps the model learn to identify objects under different lighting conditions.
- **Contrast Adjustment:** Modifying contrast allows the model to be-




Figure 13: [Left]Original [Right]Rotation

come robust against color variations, ensuring it recognizes objects regardless of their color intensity.

- **Hue Adjustment:** Random changes in hue help the model adapt to different color representations of the same object.

By adjusting these properties within the range of 0.8 to 1.2, we ensure that the changes are subtle enough to maintain the integrity of the original image.

```
def color_jitter(image, brightness=0.2, contrast=0.2, color=0.2):
    # Adjust brightness
    enhancer = ImageEnhance.Brightness(image)
    factor = random.uniform(1 - brightness, 1 + brightness)
    image = enhancer.enhance(factor)
    # Adjust contrast
    enhancer = ImageEnhance.Contrast(image)
    factor = random.uniform(1 - contrast, 1 + contrast)
    image = enhancer.enhance(factor)
    # Adjust color
    enhancer = ImageEnhance.Color(image)
    factor = random.uniform(1 - color, 1 + color)
    image = enhancer.enhance(factor)
    return image
```

Listing 8: Random Color Jittering

3.3.5 Approximate Final Size

The following is the data volume after a certain instance of data augmentation.

```
Augmented training set shape: (115017, 3072)
Augmented training set labels shape: (115017,)
```




Figure 14: [Left]Original [Right]Color Jittering

Label 0	Label 1	Label 2	Label 3	Label 4
11603	11560	11570	11470	11425
Label 5	Label 6	Label 7	Label 8	Label 9
11457	11477	11494	11494	11467

Table 2: Label Counts in training set (After augmentation)

3.4 StandardScaler

StandardScaler is used to standardize the features of the dataset by transforming them to have a mean of zero and a standard deviation of one. This normalization is crucial because it ensures that all features contribute equally to the model's performance. Without standardization, features on larger scales may dominate the learning process, leading to suboptimal model performance and slower convergence during training. By applying StandardScaler, we create a balanced dataset that enhances the model's ability to learn from all features effectively.

3.5 PCA

Principal Component Analysis (PCA) serves as a dimensionality reduction technique that helps to simplify the dataset while retaining the most significant variance. High-dimensional data can lead to overfitting and increased computational costs. This reduction helps to mitigate the curse of dimensionality, improves computational efficiency, and can also enhance the model's generalization ability.

- **Fine-tuning PCA n_components**

PCA n_components = 0.8, SVM Validation Accuracy: 0.7789

PCA n_components = 0.75, SVM Validation Accuracy: 0.7791

PCA n_components = 0.7, SVM Validation Accuracy: 0.7746

Hence, we choose 0.75 as the n_components of PCA.

3.6 SVM Classifier

From Section 3, we can notice that SVM performs significantly better than other classifiers on this dataset, so we choose not to consider adding other classifiers for ensemble learning and only use SVM.

```
svm_model = make_pipeline(  
    StandardScaler(),  
    PCA(n_components=0.75, random_state=SEED),  
    SVC(C=9, kernel='rbf', gamma='scale', random_state=SEED))  
  
svm_model.fit(X_train_featured, y_train_aug)
```

Listing 9: SVM

- **Fine-tuning C**

C=6, kernel='rbf', gamma='scale', SVM Validation Accuracy: 0.7788

C=9, kernel='rbf', gamma='scale', SVM Validation Accuracy: 0.7792

C=10, kernel='rbf', gamma='scale', SVM Validation Accuracy: 0.7792

C=12, kernel='rbf', gamma='scale', SVM Validation Accuracy: 0.7792

Hence, we choose C=9 in our SVM model.

- **Fine-tuning kernel**

C=9, kernel='rbf', gamma='scale', SVM Validation Accuracy: 0.7792

C=9, kernel='poly', gamma='scale', SVM Validation Accuracy: 0.7587

Hence, we choose kernel='rbf' in our SVM model.

4 Final Result

4.1 Performance Analysis

Using the parameters above, We finally achieve a validation accuracy of 80% and testing accuracy of 80.01% on kaggle. Our validation result is as follows:

Classification Report on Validation Set:

precision	recall	f1-score	support
-----------	--------	----------	---------




Figure 15: Kaggle Ranking

0	0.83	0.86	0.84	1011
1	0.91	0.89	0.90	1051
2	0.75	0.71	0.73	985
3	0.63	0.60	0.61	983
4	0.73	0.79	0.76	968
5	0.72	0.65	0.69	975
6	0.82	0.86	0.84	1022
7	0.87	0.84	0.86	987
8	0.87	0.88	0.88	996
9	0.85	0.91	0.88	1022
accuracy			0.80	10000
macro avg	0.80	0.80	0.80	10000
weighted avg	0.80	0.80	0.80	10000

4.2 Future Work

Based on our confusion matrix analysis, the model demonstrates strong overall performance across most categories. However, there remains room for improvement in distinguishing between label 3 and label 5, as these classes exhibit higher misclassification rates compared to others. To address this, we could implement targeted strategies in subsequent iterations, such as dedicated data augmentation techniques for these two labels (e.g., synthetic sample generation, adversarial training, or class-specific preprocessing). Alternatively, we might explore architectural adjustments like attention mechanisms or loss function modifications to enhance feature discrimination for these challenging pairs. This focused refinement could further optimize the model’s precision and robustness in real-world applications.




Figure 16: Confusion Martix